CISC-481/681 Project 2: A* Search with Heuristics

Dylan Chapp, Michael Wyatt

Department of Computer and Information Sciences

University of Delaware - Newark, DE 19716

Email: {dchapp}, {mwyatt}@udel.edu

I. INTRODUCTION

In this work we implement the A* search algorithm to solve the problem of directing a predator (e.g., a cat) through a maze of obstacles to multiple prey objectives (e.g., mice). We introduce heuristics specifically tailored to this multi-objective pathfinding problem and provide an empirical evaluation of their performance.

II. RELATED WORK

The problem presented in this report is analagous to the Traveling Sales Person (TSP) problem. In TSP, we must find the shortest route between the sales person's home, through several cities and back. It is trivial to reduce the cat and mouse problem to TSP; Therefore, the problem we are solving in this report is an NP-Complete problem, just like TSP.

The A* search algorithm has been used to solve problems like TSP, as well as many other applications. For example, A* is often used in video games for NPC map traversal. A* was developed as an extension to Dijkstra's algorithm and can be used for any general graph or tree search problems.

III. PROBLEM CHARACTERIZATION

We are presented with the problem of providing a cat with a shortest path that allows it to navigate to and consume m mice on an $n \times n$ board of tiles. However, some of these tiles are obstacles that the cat cannot traverse, which makes computing shortest paths nontrivial. We employ the A* search algorithm to find, for a fixed starting tile and a fixed goal tile containing a mouse, a shortest path for the cat. In turn, we use heuristics for executing repeated applications of single-objective A* search to find a shortest path to all of the mice.

IV. THE A* ALGORITHM

A* is an informed search algorithm used for finding the shortest path from a single node in a graph–the start–to some other node in the graph–the goal. Specifically, A* constructs a tree whose root is the start node and one of which's leaves is the goal node.

The A* algorithm bears resemblance to uninformed search algorithms such as breadth-first search in terms of its general structure, but differs in that it maintains a priority queue of nodes to expand upon rather than a FIFO/LIFO queue. The priority queue is ordered according to a heuristic that accounts for both the known cost to reach a node from the starting node and the estimated cost to reach the goal node from that node.

In fact, A* was developed as an improved upon Dijkstra's algorithm, a similiar search algorithm that too maintains a priority queue of nodes, but orders them only with respect to the known cost to reach them from the starting node. Below, 1 we provide pseudocode for A*, highlighting how A* uses the the path cost, which is used in Dijkstra's algorithm, *and* a heuristic function h to prioritize nodes in the frontier.

A. Single Objective A* Heuristics

We developed two valid heuristics for A*. Both are valid because they consistently underestimate the real cost of reaching a goal node from a given start position. Each heuristic is able to handle cases when there are multiple goals. The heuristics will help guide A* to the nearest goal on the board by returning the minimum estimated distance between a start position and any goal. The first heuristic calculates the manhattan distance between a position and goals. The manhattan distance is the number of moves–up, down, left, or right– that the cat must

Algorithm 1 A* Search

1: **function** A*(start, goal, h)
2:     explored $\leftarrow \{\}$
3:     frontier $\leftarrow$ PriorityQueue()
4:     frontier.insert(start, 0)
5:     costToReach $\leftarrow$ HashMap()
6:     parent $\leftarrow$ HashMap()
7:     costToReach[start] $= 0$
8:     parent[start] $= Nil$
9:     **while** frontier not empty **do**
10:       $current \leftarrow$ frontier.pop()
11:       **if** current == goal **then**
12:         break
13:       **else**
14:         **for** $n \in current.neighbors()$ **do**
15:           cost $\leftarrow$ costToReach[current] $+ 1$
16:           **if** $n \notin$ costToReach or
17:           cost $<$ costToReach[n] **then**
18:             costToReach[n] = cost
19:             $p = \text{cost} + h(\text{goal}, n)$
20:             frontier.insert(n, p)
21:             parent[n] = current
22:     **return** parent, costToReach

make to reach a goal. If there are no obstacles present, it is also an exact solution for the cost to a given goal. The psuedocode for this heuristic can be seen in algorithm 2.

Algorithm 2 Manhattan distance heuristic

1: **procedure** MANHATTAN(start, goals)
2:     $minDist \leftarrow \infty$
3:     $(x_1, y_1) \leftarrow start$
4:     **for** $goal \in goals$ **do**
5:       $(x_2, y_2) \leftarrow goal$
6:       $dist \leftarrow |x_2 - x_1| + |y_2 - y_1|$
7:       $minDist \leftarrow min(dist, minDist)$
    return $minDist$

The second heuristic calculates the euclidean distance between a start position and goals. The euclidean distance is the straight-line distance between any two points. This heuristic will always underestimate the distance to a goal for this problem. Pseudocode for the euclidean distance can be seen in algorithm 3.

Algorithm 3 Euclidean distance heuristic

1: **procedure** EUCLIDEAN(start, goals)
2:     $minDist \leftarrow \infty$
3:     $(x_1, y_1) \leftarrow start$
4:     **for** $goal \in goals$ **do**
5:       $(x_2, y_2) \leftarrow goal$
6:       $dist \leftarrow \sqrt{(x_2 - x_1)^2 + (y_2 - y_1)^2}$
7:       $minDist \leftarrow min(dist, minDist)$
    return $minDist$

B. Multi-Objective A* Heuristics

The A* algorithm and associated heuristics described thus far are able to handle multiple objectives, but will only find the path between a starting point and the nearest goal. If we want to contruct the path from a starting point to all goals, we need a heuristic for utilizing our single-objective A* implementation. To achieve a path from the starting point through all goals, we implemented two different multi-objective heuristics.

The first heuristic is a naïve implementation of multi-objective A*. A path to all goals is built by finding the subpath between a start point and its Nearest Neighbor (NN), which is also a goal. This is accomlished through iterative calls to the A* algorithm while changing the start position to the previously found goal. Psuedocode for this multi-objective heuristic can be found in algorithm 4.

Algorithm 4 NN multi-objective A*

1: **procedure** NEARESTGOALASTAR(start, goals)
2:     $path \leftarrow []$
3:     **while** $goals$ not empty **do**
4:       $goal, subPath \leftarrow A * (start, goals)$
5:       append $subPath$ to $path$
6:       $start \leftarrow goal$
7:       remove $goal$ from $goals$
    return $path$

The second heuristic works by exhaustively searching all possible paths and returning the path which touches each goal and has the lowest cost. This path can thought of as a Minimum Spanning Path (MSP) for the start and goals. The algorithm first builds a library of all possible subpaths between any two points in the set of start and goal positions. Using this library, it then builds all possible paths

of an appropriate length ($len(goals) - 1$). The algorithm then removes any path which does not visit all goals and the start position (i.e., a complete path.) A path cost for each of these paths is then computed and the path with the minimum cost is returned. Pseudocode for this multi-objective heuristic can be found in algorithm 5.

For the pseudocode in algorithm 5, the function $CompletePaths(subPaths)$ builds all possible paths of length $length(goals) - 1$ and returns only the complete paths. The fuction $GetPathCost(path, cost)$ calculates the cost of a path using the path and a list of costs for all subpaths.

Algorithm 5 MSP multi-objective A*

1: **procedure** COMBINATORIALASTAR(start, goals)
2:     $cost \leftarrow []$
3:     $subPaths \leftarrow$ combinatorial pairs of $\{start, goals\}$
4:     **for** $subPath \in subPaths$ **do**
5:         $(start, goal) \leftarrow subPath$
6:         $pathCost \leftarrow A*(start, goal)$
7:         append $pathCost$ to $cost$
8:     $paths \leftarrow CompletePaths(subPaths)$
9:     $minCost \leftarrow \infty$
10:     $minPath \leftarrow NULL$
11:     **for** $path \in paths$ **do**
12:         $pathCost \leftarrow GetPathCost(path, cost)$
13:         **if** $pathCost < minCost$ **then**
14:             $minPath \leftarrow path$
15:             $minCost \leftarrow cost$
        return $minPath$

V. IMPLEMENTATION

We implemented our core single-objective A* search algorithm, all helper functions for ingesting and reporting data, and all functions necessary to adapt the single-objective algorithm to the multi-objective problem posed in Python. This choice was guided by the authors' familiarity with the language and the knowledge that we would be able to rapidly iterate on and test our design. With the exception of a few uses of the NumPy numerical computation module for convenience, the project is implemented without the use of 3rd-party modules.

A. Components

Our implementation is composed of three primary components: board generation, the core single-objective solver and its heuristics, and a top-level driver program that invokes the single-objective solver according to its own set of heuristics to solve multi-objective boards.

In addition to the board specified in the project description, our implementation can read in an arbitrary board from a file or generate a broader class of randomized board from user-specified parameters. The initial impetus for providing a more flexible board-generation procedure than that specified in the project description was to explore the effect of the distribution of obstacles on the board on runtime for a fixed board size.

The single-objective solver essentially implements A* as specified by the pseudocode listed above. Though it would be possible to implement the algorithm purely in terms of references to positions of the board–e.g., for enumerating neighbors of a square– we generate a graph representation of the board prior to path-finding. Not only does this lead to a cleaner implementation of the core algorithm, but it also allows the solver to generalize to boards of dimension than 2, simply by redefining the helper function for enumerating neighbors.

B. Use

To use our implementation, first ensure that all dependencies are met as per the README available at https://github.com/dchapp/cisc681_search_project and then invoke the `Driver.py` program as specified in the README. By default, the program will generate and solve a 10×10 board conforming to the problem statement's obstacle and objective frequencies. Also by default, the program will use a fixed choice of both single-objective and multi-objective heuristic. However, the size of board, board-generation procedure, and choice of heuristic can all be given as arguments to the `Driver.py` program.

VI. EVALUATION

There are two factors upon which we judge the quality of our implementation: runtime performance and optimality of the generated path

A. Runtime Performance

Each combination of a single-objective heuristic and a multi-objective heuristic–e.g., the Manhattan distance single-objective heuristic and the nearest-neighbor multi-objective heuristic–is executed 20 times on boards of increasing size for a fixed number of objectives. In initial testing we observed large variation in runtime due to differences in obstacle layout on the randomly generated boards. Consequently, we report the mean and standard deviation of the runtimes for each heuristic combination.


Fig. 1. Mean of runtimes for variable-width boards with a fixed number of objectives


Fig. 2. Standard deviation of runtimes for variable-width boards with a fixed number of objectives

The time complexity of the single-objective A* algorithm is $O(b^d)$ where d is the depth of the optimal path to the nearest goal and b is the average branching factor. In practice, we assume $b = 3$ since for any node, there are three neighbors that can be expanded. For the multi-objective procedures that invoke single-objective A*, the nearest-neighbor strategy has time complexity $O(g)$ and the minimum spanning-path strategy has time compl exity $O(g^2)$ where g is the number of goals and we assume an average constant cost of calls to single-objective A* by amortizing over the goals.

We observe agreement between the empirical runtime data and the expected theoretical perf ormance of the algorithms-heuristic combinations. However, in future work we would like to explore larger boards and vary the number of obj ectives to put this claim on firmer footing.

B. Optimality

The A* algorithm is guaranteed to return the shortest path from a start to end objective if the heuristic used in admissible. Both the manhattan and euclidean heuristic we developed are admissible for this search problem. By using the nearest goal, we are guaranteed to find the shortest path to *any* goal. Additionally, the euclidean distance and manhattan distance will always be less than the or equal to the actual cost of reaching a goal node. This is true for Euclidean distance because it calculates the straight line distance between a start and goal node. The straight line between any two points is the shortest path between those points. The manhattan distance will also underestimate the cost of reaching a goal because it calculates distance based on the same rules that the cat ues to move (i.e., can only move horizontally or vertically.) The manhattan heuristic also ignores obstacles, so it will be less than or equal to the actual distance between a start and goal.

The Nearest Neighbor multi-objective heuristic for this search problem is not always optimal. However, it often produces paths which are very close to optimal. The reason for non-optimality is that the algorithm will always choose the path to the nearest goal to the current position. This does not account for the distribution of the remaining goals on the map. As an example, figure 3 shows an input map and the non-optimal solution found by the Nearest Neighbor heuristic.

The Minimum Spanning Path multi-objective heuristic will always find the optimal path, but at
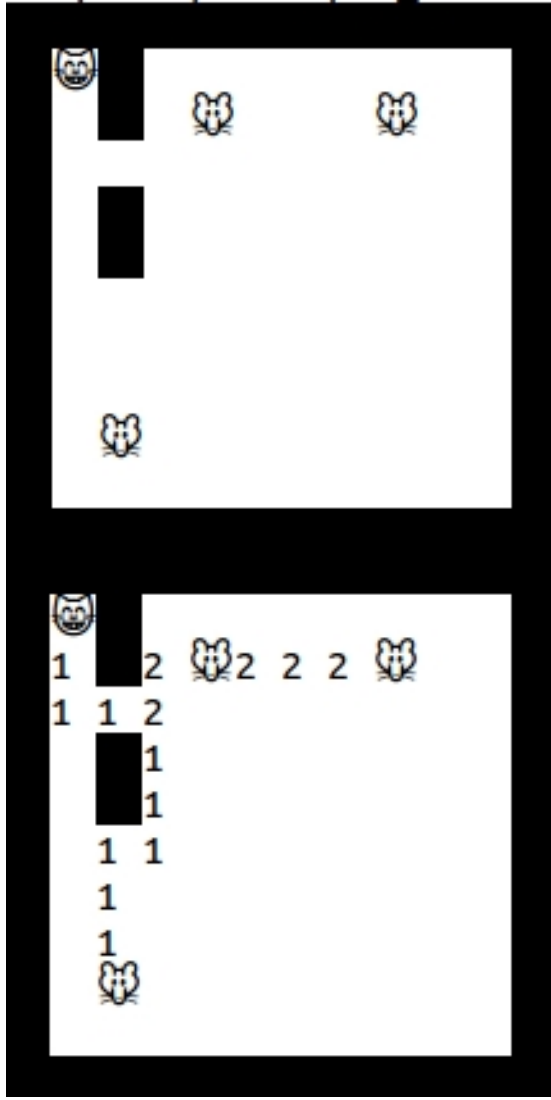
Fig. 3. An example multi-objective map and the Nearest Neighbor heuristic's non-optimal solution. Path length: 23

Fig. 4. The optimal solution found from the Minimum Spanning Path heuristic. Path length: 22

the cost of time complexity. Because this heuristic computes the cost of all possible paths through the start point and goals, it takes much longer to compute with large boards. However, it will always produce the most optimal path. The optimal path determined by this heuristic for the same map as in figure 3 can be seen in figure 4.

VII. FUTURE WORK

The work we have presented in this report is complete for the purposes of cat-and-mouse problem. However, as the size and complexity of the problem becomes larger, our methods will need to be expanded and optimized to allow for faster com-

putation of a shortest path. In order to accomodate very large boards and high dimension boards, the combinatorial method of the MSP heuristic will become increasingly slow. Optimizations can be made by reducing the number of paths which are tested. For example, rather than testing every possible path, some of which may not be connected, build and test only paths which touch every goal. Additionally, parallelizing the multi-objective heuristics with MPI or OpenMP would give a substantial increase in performance.

VIII. CONCLUSION

We have explored the impact of heuristic choice on a multi-objective search problem utili zing the A* search algorithm. Multiple heuristics for both single-objective A* and the multi-objective routine utlizing A* are examined theoretically and tested for runtime performance. In addition to our implementation of A* and the heuristics, our application provides diverse methods for constructing instances of the search problem and allows future developers to extend the solvers we provide by adding their own heuristics and multi-objective strategies.